

# Prior SBC on the Lotka-Volterra model

Teemu Säilynoja

2024-06-20

In this notebook, we show the results of calibration assessment on the [Lotka-Volterra model](#) using prior SBC. For collected results of all our experiments, see the notebook named [lotka-volterra-sbc](#).

## 1 Setup

Like the original Hudson Bay data, we simulate data sets of 21 years. We run the SBC on 250 prior predictive samples, each with a total of 4000 post warm-up posterior samples, what will be thinned down by a factor of 10, yielding 400 posterior draws for calculating the rank of the prior draws and joint log-likelihood.

## 2 The Model

```
functions {
  vector dz_dt(real t,          // time
               vector z,        // system state {prey, predator}
               array[] real theta // parameters
               //real[] x_r,    // unused data
               //int[] x_i
               ) {
    real u = z[1];
    real v = z[2];

    real alpha = theta[1];
    real beta = theta[2];
    real gamma = theta[3];
    real delta = theta[4];
```

```

    real du_dt = (alpha - beta * v) * u;
    real dv_dt = (-gamma + delta * u) * v;

    return to_vector({du_dt, dv_dt});
  }
}

data {
  int<lower = 0> N;          // number of measurement times excl. year 0
  array[N] real ts;         // measurement times > 0
  array[2] real y_init;     // initial measured populations
  array[N,2] real <lower = 0> y; // measured populations
}

parameters {
  array[4] real<lower = 0> theta; // { alpha, beta, gamma, delta }
  vector<lower = 0> [2] z_init; // initial population
  array[2] real<lower = 0> sigma; // measurement errors
}

transformed parameters {
  array[N] vector[2] z
    = ode_rk45_tol(dz_dt, z_init, 1.0, ts, 1e-6, 1e-5, 1000, theta);
}

model {
  theta[{1, 3}] ~ normal(1, 0.5);
  theta[{2, 4}] ~ normal(0.05, 0.05);
  sigma ~ lognormal(-1, 1);
  z_init ~ lognormal(log(10), 1);
  for (k in 1:2) {
    y_init[k] ~ lognormal(log(z_init[k]), sigma[k]);
    y[, k] ~ lognormal(log(z[, k]), sigma[k]);
  }
}

generated quantities {
  real loglik = 0;
  vector[N + 1] log_lik;
  log_lik[1] = 0;
  for (k in 1:2) {
    log_lik[1] += lognormal_lpdf(y_init[k]|log(z_init[k]), sigma[k]);
    loglik += lognormal_lpdf(y_init[k]|log(z_init[k]), sigma[k]);
    loglik += lognormal_lpdf(y[, k]|log(z[, k]), sigma[k]);
  }
}

```

```

    }
    for (n in 1:N) {
      log_lik[n + 1] = 0;
      for (k in 1:2) {
        log_lik[n + 1] += lognormal_lpdf(y[n, k] | log(z[n, k]), sigma[k]);
      }
    }
  }
}

```

## 3 Simulation-based calibration

### 3.1 Prior predictive samples

For implementing the SBC, we use the [SBC R-package](#) by Angie H. Moon et al. First, we define the data generating process, which in our case is generating prior predictive samples.

Next, we use the data generating process to define a data generator and construct 250 data sets for SBC.

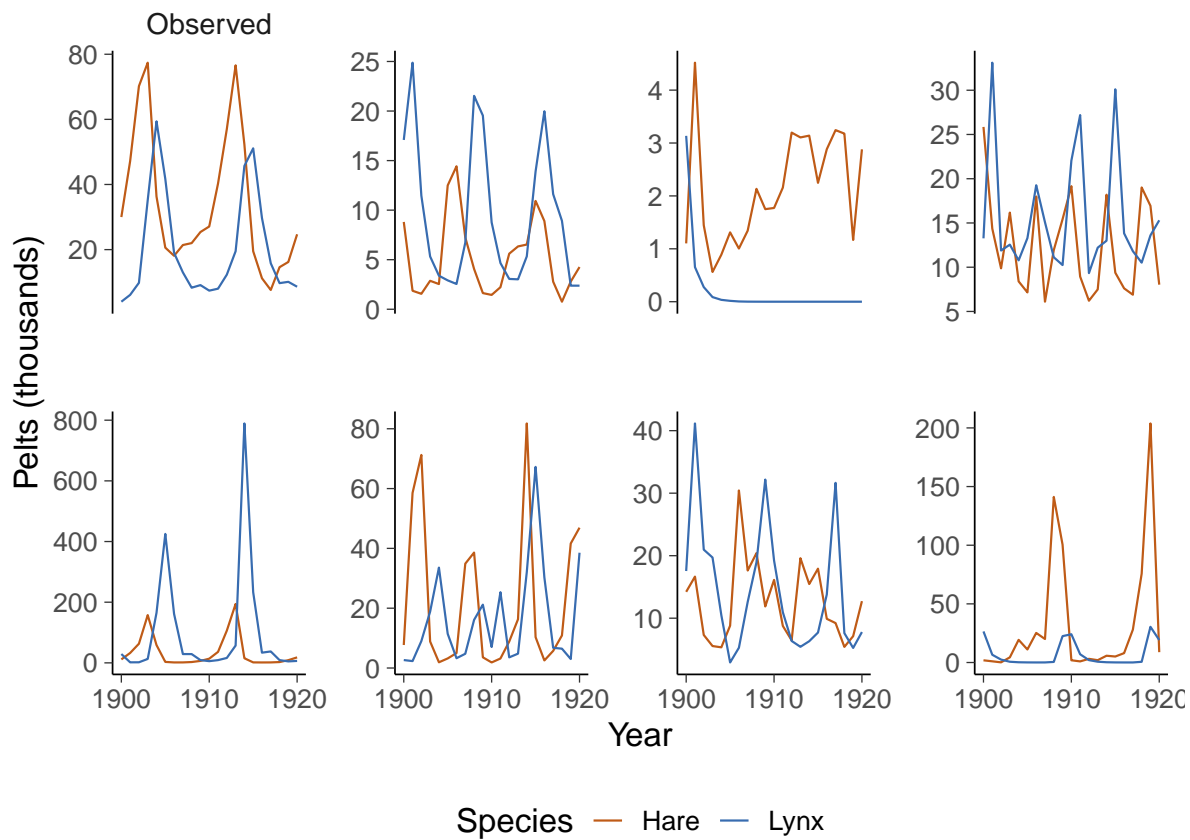


Figure 1: Observed data compared to prior predictive draws.

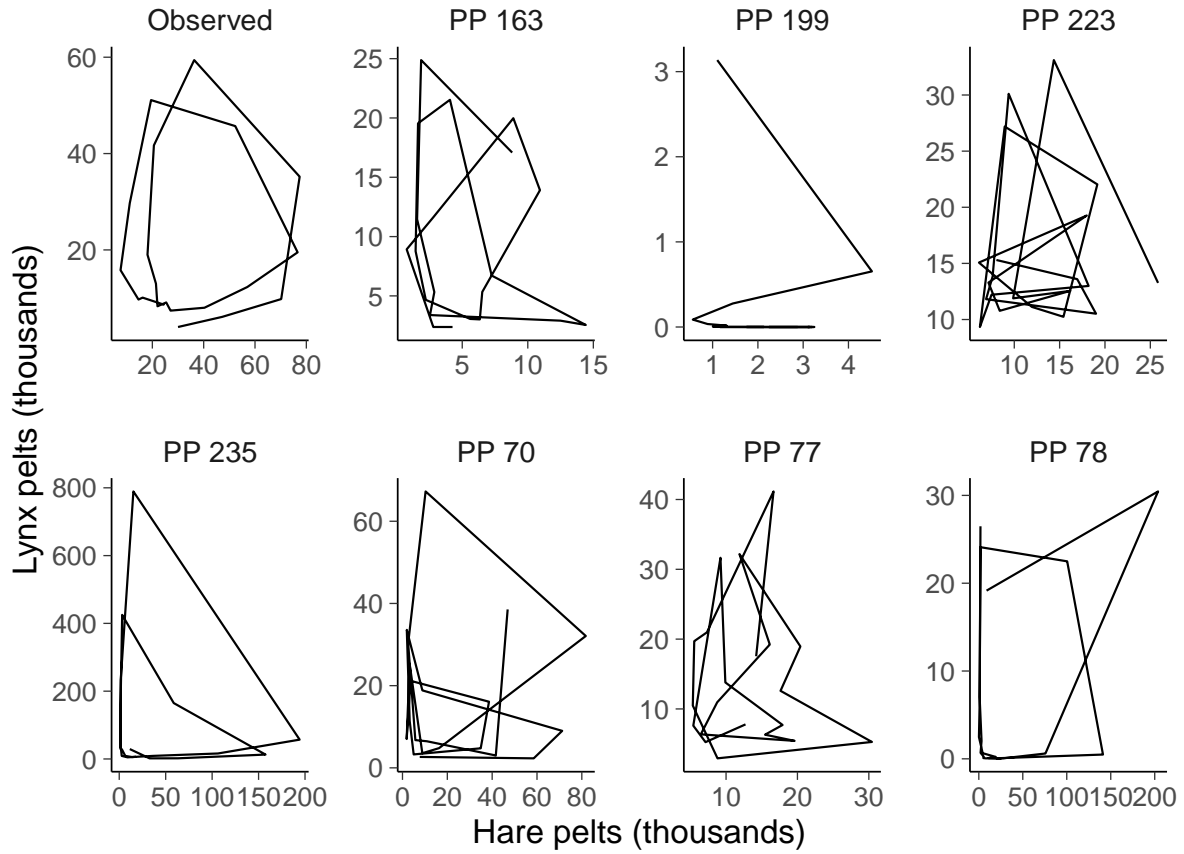


Figure 2: Trajectories of the population dynamics in the observation and the prior predictive draws from above.

### 3.2 Posterior samples

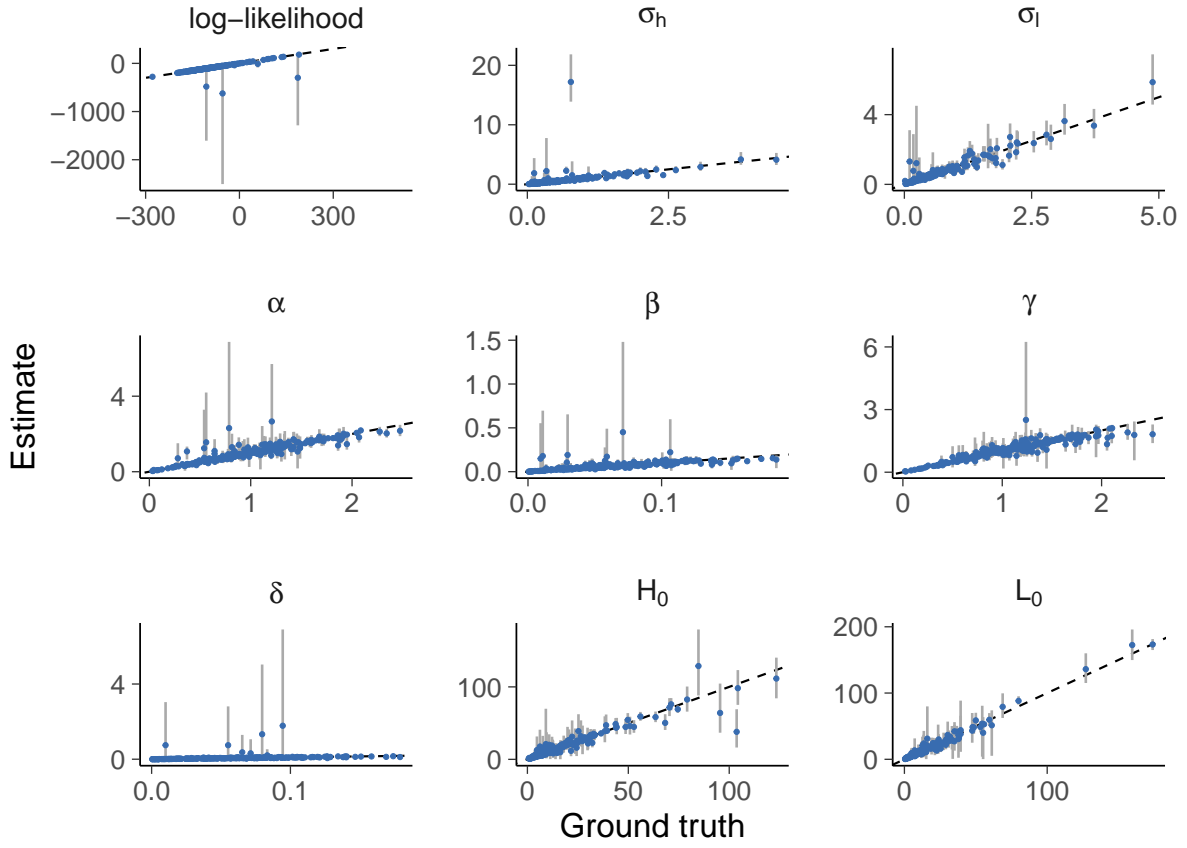
We use `cmdstanr` to obtain posterior draws via MCMC sampling.

### 3.3 Run SBC

This is the step requiring a lot of computation. We run MCMC on each of the predictive samples and store statistics of the results.

### 3.4 Results

A simple plot to compare the posterior samples and the ground truth values.



The ECDF of the PIT values of the prior draw with regards to the posterior sample gives a principled way to check for the calibration of the model, as we can draw simultaneous 95% confidence intervals for the PIT-ECDFs. The joint log-likelihood is a good test quantity for overall calibration of the inference, and shows here some possible calibration issues, although the number of SBC iterations is quite low.

